

# Week 5

## Relational Databases

**The Relational Model, Relational Algebra, and SQL**

# Today's Topics

- **Part I:**
  - Introduction to the **relational model** of database design
  - Introduction to **relational algebra**, which forms the basis of the **Structured Query Language (SQL)**
  - Mapping relational algebra to SQL
- **Part II:**
  - **Hands-on exercise** with SQL

# Database Options

Repeated from Week 4

## SQL databases (or relational databases)

- Well defined **schema**
  - the database is modeled as a set of joinable tables with pre-defined columns
- Entities are stored in **tables** (called relations), managed by an **RDBMS**
- **Tables can be joined** and using a powerful and flexible querying language: **SQL**
- Data follows a **rigid structure** and is mostly **non-redundant**
- Changing/updating the structure of data involves change in schema and **database alterations**.
- There is a DBA – a **database administrator** -- between the application developer and the database, responsible for defining and altering the schema if needed.
  - Using **SQL's DDL**

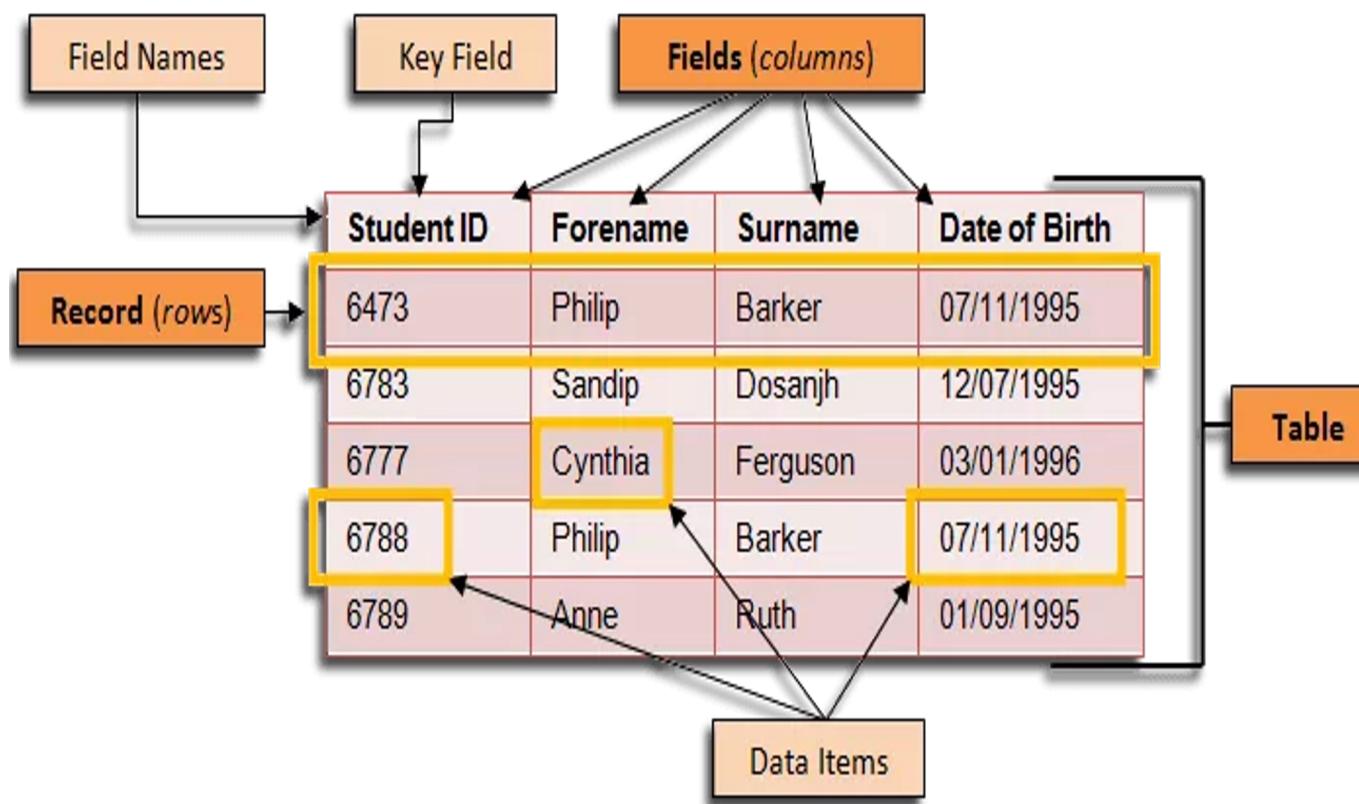
**Naturally ideal for vertical scaling, i.e.,** increasing the database's capacity by adding more resources (CPU, memory, storage) to a single machine.

# The Relational Model

A **relation** (table) contains **records** (rows) made of **fields** (columns).

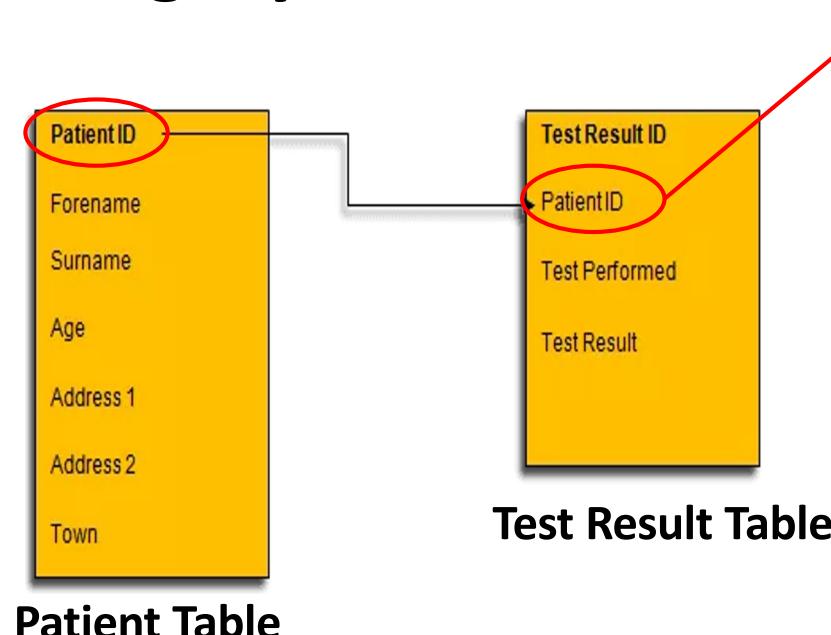
At least one of the fields is a **primary key** (key field), used to uniquely identify a record.

**Important rule:** each data item has a **type** (the type of that field) and must be atomic, i.e., it cannot be a list or any other kind of collection.



# The Relational Model

**Foreign Key:** a field (or set of fields) in one relation that references the primary key of another relation, establishing a relationship and enforcing **referential integrity** between the two relations.



Foreign Key

**Referential integrity** is enforced, as only those patient IDs that exist in the Patient table could appear as foreign keys in the Test Result table

Referential Integrity is enforced by an RDBMS when, for example:

- A patient is removed from the Patient table.
- A test result is added to the Test Result table.

# Relational Schema

- The name of a relation and its fields together is called the **schema** of that relation.

e.g. the schema for a Movies relation may be written as:

**Movies (title, year, length, genre)**

- The Movies table contains rows each of which have 4 columns, namely: title, year, length and genre (usually displayed in that order).
- title and year together are the **primary key (composite)**, i.e. the (title, year) combination uniquely identifies a row of the Movies table

- The collective schema of all relations of a database make up the **database schema**.
- **In practice**, to create a relational database, a database administrator must specify the database schema using SQL's DDL (data definition language).

# Example: A simple HR database

**Employee** (fname, mint, lname,  
empid, bdate, address, gender,  
salary, superid, dno)

**Department** (dname, dno, mgrid,  
mgstartdate)

**Project** (pname, pno, plocation,  
dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname,  
gender, bdate, relationship)

## Overall description:

Employees work for departments.  
Departments run projects.  
Employees work on projects.  
Employees have dependents.

## Foreign keys:

Employee.dno  
Employee.superid  
Department.mgrid  
Project.dno  
WorksOn.empid  
WorksOn.pno  
Dependent.empid

# Relational Operators

- **Relational operators** operate on existing relations to produce new relations corresponding to user queries
- Relational operators maybe divided into two main categories:

Set Operators	Database Operators
Union Intersection Difference Cartesian Product	Select Project Join Aggregate Division Rename

# Projection

Extracts certain columns from the table:

$$Temp1 = \prod_A(r)$$

r	A	B	C
1	610	3	
1	620	3	
1	600	2	
1	650	2	
2	610	3	
2	634	4	

$$Temp2 = \prod_{B,C}(r)$$

Temp1	A
1	
2	

Temp2	B	C
610	3	
620	3	
600	2	
650	2	
634	4	

Set Semantics

Corresponding SQL:

SELECT A FROM r;

SELECT B, C FROM s;

Bag Semantics (requires DISTINCT for set semantics)

# Union

- Written as  $(r+s)$  or  $(r \cup s)$
- Relations must be union compatible
- Duplicate rows are eliminated

r	A	B	C
1	1	1	1
2	2	2	2
3	3	3	3

s	A	B	C
1	1	2	3
1	1	1	1
3	2	1	

r+s	A	B	C
1	1	1	1
2	2	2	2
3	3	3	3
1	2	3	
3	2	1	

Corresponding SQL:  
SELECT \* FROM r  
UNION  
SELECT \* FROM s;

# Intersection

- Written as  $(r \cap s)$
- Relations must be intersection compatible

r	A	B	C
1	1	1	1
2	2	2	2
3	3	3	3

s	A	B	C
1	2	3	
1	1	1	1
3	2	1	

$r \cap s$	A	B	C
1	1	1	1

Corresponding SQL:  
SELECT \* FROM r  
INTERSECT  
SELECT \* FROM s;

# Difference

- Written as  $(r-s)$
- Relations have to be difference compatible
- Includes rows that are in  $r$  but not in  $s$

r	A	B	C
1	1	1	1
2	2	2	2
3	3	3	3

s	A	B	C
1	1	2	3
1	1	1	1
3	2	1	

r-s	A	B	C
2	2	2	2
3	3	3	3

Corresponding SQL:  
SELECT \* FROM r  
EXCEPT  
SELECT \* FROM s;

# Queries Using Union, Intersection, Difference and Projection

*List the ids of all employees who are working on some project.*

$$\text{ans} = \prod_{\text{empid}} (\text{WorksOn})$$

```
SELECT empid  
FROM WorksOn;
```

*List the ids of all employees who are not working on any project.*

$$\text{ans} = \prod_{\text{empid}} (\text{Employee}) - \prod_{\text{empid}} (\text{WorksOn})$$

```
SELECT empid FROM Employee  
EXCEPT  
SELECT empid FROM WorksOn;
```

*List the ids of all employees who are working on a project and have a dependent.*

$$\text{ans} = \prod_{\text{empid}} (\text{WorksOn}) \cap \prod_{\text{empid}} (\text{Dependent})$$

```
SELECT empid FROM WorksOn  
INTERSECT  
SELECT empid FROM Dependent;
```

## Schema

**Employee** (fname, mint, lname, empid, bdate, address, gender, salary, superid, dno)

**Department** (dname, dno, mgrid, mgstartdate)

**Project** (pname, pno, plocation, dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender, bdate, relationship)

# Queries Using Union, Intersection, Difference and Projection

*List the ids of all supervisors who either have dependents or are working on a project or both.*

$$t1 = \prod_{empid} (\text{WorksOn}) \cap \prod_{superid} (\text{Employee})$$

these are the supervisors who are working on a project

$$t2 = \prod_{superid} (\text{Employee}) \cap \prod_{empid} (\text{Dependent})$$

these are the supervisors who have dependents

$$ans = t1 \cup t2$$

## Schema

**Employee** (fname, mint, lname, empid, bdate, address, gender, salary, superid, dno)

**Department** (dname, dno, mgrid, mgstartdate)

**Project** (pname, pno, plocation, dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender, bdate, relationship)

```
SELECT empid FROM WorksOn  
INTERSECT  
SELECT superid FROM Employee  
  
UNION  
  
SELECT superid FROM Employee  
INTERSECT  
SELECT empid FROM Dependent;
```

**Corresponding SQL**

# Conditional Selection

The selection operator extracts certain rows from the table and discards the others. Retrieved tuples must satisfy a given **filtering condition**.

$$Temp1 = \sigma_{(B \geq 620) \text{ and } (C < 4)} (r)$$

r	A	B	C
1	610	3	
1	620	3	
1	600	2	
1	650	2	
2	610	3	
2	634	4	

Temp 1	A	B	C
1	620	3	
1	650	2	

```
SELECT *
FROM r
WHERE B >= 620 AND C < 4;
```

Corresponding SQL

# Queries Using Union, Intersection, Difference, Projection and Conditional Selection

*List the names and genders of all employees with a salary of at least 80, 000 pounds.*

$$ans = \prod_{\text{fname, mint, lname, gender}} (\sigma_{\text{salary} \geq 80000}(\text{Employee}))$$

```
SELECT fname, mint, lname, gender  
FROM Employee  
WHERE salary >= 80000;
```

*List the ids of all employees with a salary of at least 80, 000 pounds who are not working on any project.*

$$ans = \prod_{\text{empid}} (\sigma_{\text{salary} \geq 80000}(\text{Employee})) - \prod_{\text{empid}} (\text{WorksOn})$$

## Schema

**Employee** (fname, mint, lname, empid, bdate, address, gender, salary, superid, dno)

**Department** (dname, dno, mgrid, mgstartdate)

**Project** (pname, pno, plocation, dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender, bdate, relationship)

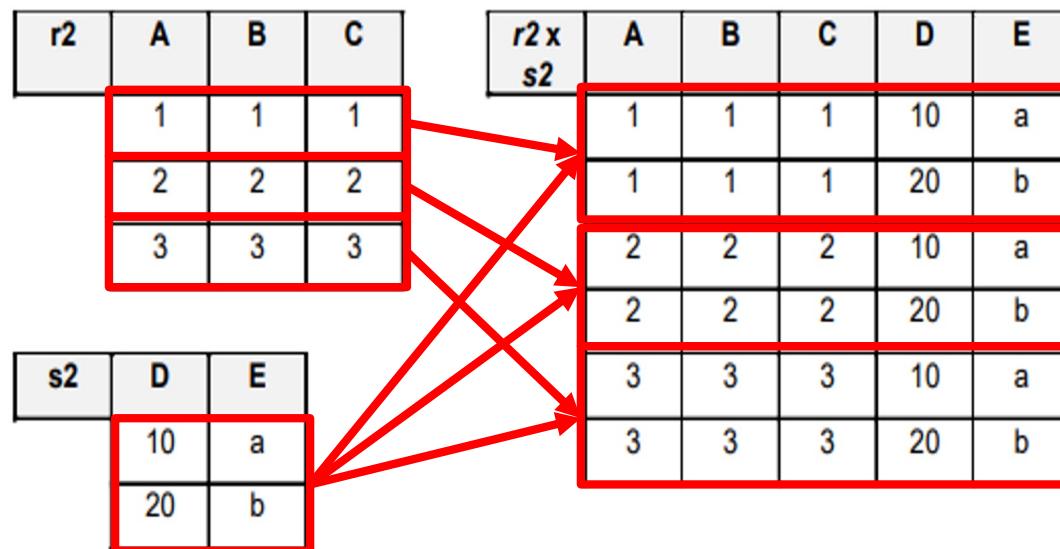
```
SELECT empid  
FROM EMPLOYEE  
WHERE salary >= 80000
```

EXCEPT

```
SELECT empid  
FROM WorksOn
```

# Cross Product (or Cross Join)

- Written as  $(r \times s)$
- Concatenates rows from two relations, making all possible combinations of rows.



```
SELECT *  
FROM r2, s2;
```

Corresponding SQL

# Join (or inner join)

- The join operation, denoted by  $(r \bowtie_{\text{COND}} s)$ , is used to combine *related tuples* from two relations.
- Here COND is the matching condition  $r \bowtie_{\text{COND}} s$
- The following example demonstrates the operation  $r \bowtie_{C=D} s$

r	A	B	C
1	1	1	1
2	2	2	2
3	3	3	3

s	D	E
1	a	
2	b	
2		c

Temp1	A	B	C	D	E
1	1	1	1	1	a
2	2	2	2	2	b
2	2	2	2	2	c

```
SELECT *
FROM r JOIN s
ON r.C = s.D;
```

Corresponding SQL

# Natural Join

Denoted by  $(r * s)$

Combines tuples of two relations using an implicit condition, i.e. the tables are related by columns that have the same names and types.

r	A	B	C
1	1	1	
2	1	0	
4	3	2	

s	B	C	D
1	1	a	
1	2	b	
3	2	c	
4	3	d	

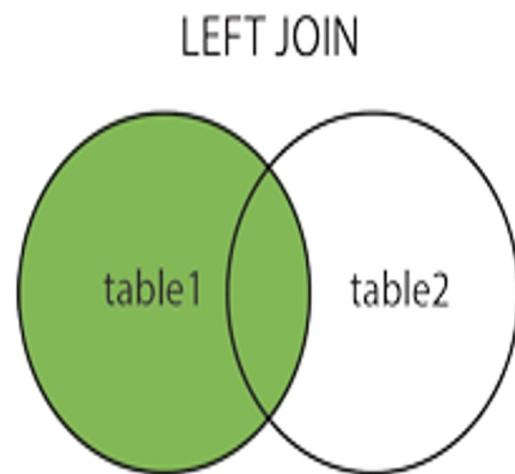
Temp	A	B	C	D
1	1	1	1	a
4	3	2	c	

SELECT \*  
FROM r NATURAL JOIN s;

Corresponding SQL

# Left Outer Join

- Denoted  $(r \bowtie_{\text{COND}} s)$
- Combines tuples of two relations and keeps in the result every tuple from the left table, but only those from the right table that meet the join condition.



```
SELECT *
FROM r LEFT JOIN s
ON r.A = s.D;
```

Corresponding SQL

# Left Outer Join Example

( $r \bowtie_{A=D} s$ )

r	A	B	C
1	1	1	
2	2	2	
3	3	3	

s	D	E
1	a	
2	b	
2	c	

Temp1	A	B	C	D	E
1	1	1	1	1	a
2	2	2	2	2	b
2	2	2	2	2	c
3	3	3	3	null	null

```
SELECT *
FROM r LEFT JOIN s
ON r.A = s.D;
```

Corresponding SQL

# Queries Using Set Operations, Projection, Conditional Selection and Joins

*List the names, genders and department names of all employees.*

$$ans = \prod_{fname, mint, lname, gender, dname} (\text{Employee} * \text{Department})$$

```
SELECT fname, mint, lname, gender, dname  
FROM Employee NATURAL JOIN Department;
```

*List the names of all the female employees in the HR department.*

$$ans = \prod_{fname, mint, lname} (\sigma_{\text{gender}=\text{"F"}}(\text{Employee}) * \sigma_{\text{dname}=\text{"HR"} }(\text{Department}))$$

alternatively,

$$ans = \prod_{fname, mint, lname} (\sigma_{\text{gender}=\text{"F"} \text{ and } \text{dname}=\text{"HR"} } (\text{Employee} * \text{Department}))$$

Which of the two alternatives is better?

## Schema

**Employee** (fname, mint, lname, empid, bdate, address, gender, salary, superid, dno)

**Department** (dname, dno, mgrid, mgstartdate)

**Project** (pname, pno, plocation, dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender, bdate, relationship)

```
SELECT fname, mint, lname  
FROM Employee e JOIN Department d  
ON e.dno = d.dno  
WHERE e.gender = 'F'  
AND d.dname = 'HR';
```

RDBS performs query chain optimization to generate an optimal sequence of relational algebra statements.

# Queries Using Set Operations, Projection, Conditional Selection and Joins

*List the ids of all employees of the EEE department working on project Panopto.*

ans

$$= \prod_{\text{empid}} (\text{Employee} * (\sigma_{\text{dname}=\text{"EEE"}}(\text{Department}) * \sigma_{\text{pname}=\text{"Panopto"}}(\text{Project})))$$

which attribute is being used to join?

```
SELECT DISTINCT e.empid
FROM
Employee e JOIN Department d
ON e.dno = d.dno
JOIN WorksOn w
ON e.empid = w.empid
JOIN Project p
ON w.pno = p.pno
WHERE d.dname = 'EEE'
AND p.pname = 'Panopto';
```

## Schema

**Employee** (fname, mint, lname, empid, bdate, address, gender, salary, superid, dno)

**Department** (dname, dno, mgrid, mgstartdate)

**Project** (pname, pno, plocation, dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender, bdate, relationship)

*List the last names of all employees who have at least one dependent.*

ans =  $\prod_{\text{empid}}$  Dependent  
SELECT empid FROM Dependent;

*Get the last names of all employees who have at least two dependents.*

???

# Aggregation

- Syntax:  $(_{\text{<Grouping attribute>}} F_{\text{<function list>}} (\text{relation name}) )$
- **<function list>** : contains simple mathematical functions.
  - Common functions are: MAX, MIN, AVG, SUM and COUNT
- **<Grouping attribute>**: contains a column name to organize the function outputs into groups.
  - Without a grouping attribute a global, across-the-table result is returned

# Aggregation Example

$$(\underset{A}{\underset{\text{F}}{\underset{\text{SUM(B), MAX(C)}}{\text{(r)}}})}$$

r	A	B	C
1	10	1	
1	2	5	
2	3	3	
3	6	10	
3	5	7	

Temp	Group-by field	Summary Data	
	A	Sum_B	Max_C
1	10	12	5
2	3	3	3
3	6	11	10

```
SELECT A,  
       SUM(B) AS Sum_B,  
       MAX(C) AS Max_C  
FROM r  
GROUP BY A;
```

Corresponding SQL

Renaming of column names required to have reusable column names, e.g., to be used in projection or conditional selection, etc.

# Queries Using Aggregation

*Get the employee ids and number of dependents of all employees who have at least two dependents.*

```
SELECT e.empid,  
       COUNT(*) AS num_dependents  
  FROM  
Employee e JOIN Dependent d  
    ON e.empid = d.empid  
  
 GROUP BY e.empid  
  
 HAVING COUNT(*) >= 2;
```

## Schema

**Employee** (fname, mint, lname, empid, bdate, address, gender, salary, superid, dno)

**Department** (dname, dno, mgrid, mgstartdate)

**Project** (pname, pno, plocation, dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender, bdate, relationship)

**HAVING** is used for conditional selection when  
**GROUP BY** based aggregation is involved.

# Nested Queries (or Subqueries)

*List the ids and names of employees who are working on some project.*

```
SELECT empid, fname, mint, lname  
FROM Employee  
WHERE empid IN (  
    SELECT empid  
    FROM WorksOn  
)
```

*List the ids and names of employees who have at least one dependent.*

```
SELECT e.empid, fname, mint, lname  
FROM Employee e  
WHERE EXISTS (  
    SELECT *  
    FROM Dependent d  
    WHERE d.empid = e.empid  
)
```

## Schema

**Employee** (fname, mint, lname, empid, bdate, address, gender, salary, superid, dno)

**Department** (dname, dno, mgrid, mgstartdate)

**Project** (pname, pno, plocation, dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender, bdate, relationship)

## Other Subquery Operators like IN that can be used:

**IN**: value appears in the subquery result set

**NOT IN**: value does not appear in the subquery result set

**EXISTS**: subquery returns at least one row

**NOT EXISTS**: subquery returns no rows

**ANY**: comparison true for at least one value returned by subquery

**SOME**: same as ANY

**ALL**: comparison true for every value returned by subquery

# Miscellaneous Queries

*List the complete employee data of all employees, and if they are working on projects, display the total number of hours they spend across all their projects.*

```
SELECT e.*, SUM(w.hours) AS total_hours  
  
FROM Employee e LEFT JOIN WorksOn w  
ON e.empid = w.empid  
  
GROUP BY e.empid;
```

## Schema

**Employee** (fname, mint, lname, empid, bdate, address, gender, salary, superid, dno)

**Department** (dname, dno, mgrid, mgstartdate)

**Project** (pname, pno, plocation, dno)

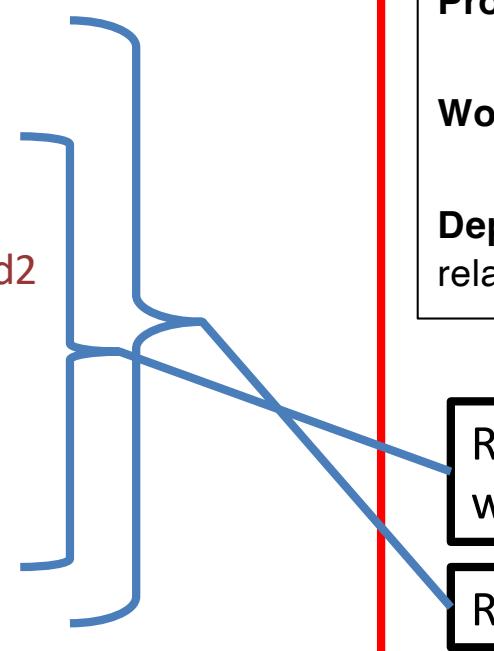
**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender, bdate, relationship)

# Miscellaneous Queries

*List the ids and last names of all EEE employees working on maximum number of projects.*

```
SELECT e.empid, e.lname  
FROM Employee e  
JOIN Department d ON e.dno = d.dno  
JOIN WorksOn w ON e.empid = w.empid  
WHERE d.dname = 'EEE'  
GROUP BY e.empid, e.lname  
HAVING COUNT(*) = (  
    SELECT MAX(ct)  
    FROM (  
        SELECT COUNT(*) AS ct  
        FROM  
        Employee e2 JOIN Department d2  
        ON e2.dno = d2.dno  
        JOIN WorksOn w2  
        ON e2.empid = w2.empid  
        WHERE d2.dname = 'EEE'  
        GROUP BY e2.empid  
    )  
);
```



## Schema

**Employee** (fname, mint, lname, empid, bdate, address, gender, salary, superid, dno)

**Department** (dname, dno, mgrid, mgstartdate)

**Project** (pname, pno, plocation, dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender, bdate, relationship)

Results in a single-column table with column name ct

Results in a single value.

# Miscellaneous Queries

*Output a table with the following columns:*

*Employee Id, Last Name, Dept. name, Salary, No. of Proj,  
No. of Dependents*

```
SELECT
    e.empid AS "Employee Id",
    e.lname AS "Last Name",
    d.dname AS "Dept. name",
    e.salary AS "Salary",
    COUNT(DISTINCT w.pno) AS "No. of Proj",
    COUNT(DISTINCT dep.depname) AS "No. of Dependents"
FROM Employee e JOIN Department d
ON e.dno = d.dno
JOIN WorksOn w
ON e.empid = w.empid
JOIN Dependent dep
ON e.empid = dep.empid
GROUP BY
    e.empid, e.lname, d.dname, e.salary;
```

## Schema

**Employee** (fname, mint, lname,  
empid, bdate, address, gender, salary,  
superid, dno)

**Department** (dname, dno, mgrid,  
mgstartdate)

**Project** (pname, pno, plocation, dno)

**WorksOn** (empid, pno, hours)

**Dependent**(empid, depname, gender,  
bdate, relationship)

# RDBMS

A **database** is a structured collection of data that is managed by a DBMS (Database management system)

- Some of the most powerful and widely used database management systems are relational: **RDBMS**.
  - They provide a **high-level query language** called SQL, which is based on relational algebra.
  - Examples of RDBMS: Oracle, MySQL, Microsoft SQL Server, PostgreSQL, **SQLite**, Amazon RDS, etc.

# RDBMs Ensure ACID Properties

- **Atomicity:** transactions, which consist of multiple statements, are treated as atomic.
  - Either the whole thing succeeds or the whole thing fails
- **Consistency:** transactions can only bring the database from one valid state to another
  - a state is valid if it follows all the specified rules including constraints
- **Isolation:** ensures that concurrent execution of transactions leaves the database in the same state that would have been obtained if the transactions were executed sequentially.
- **Durability:** database commits are final and durable even if system fails after a commit.

# PART II

## Practicing with SQL

- **Open:**  
<https://sqliteonline.com/>
- **Upload Schema/Data:**
  - Download [emp\\_db.sql](#)
  - Import into <https://sqliteonline.com/>

# Write SQL for the Following Queries

- 1) List the first name, last name, and salary of all employees whose salary is at least 80000.
- 2) List the employee id and department number of all employees whose gender is 'F' and who live at an address containing 'London' (use LIKE '%London%')
- 3) List the department name and manager start date for all departments whose department number is greater than 20.

**Expected result:**

dname	mgstartdate
Finance	2021-04-10
Sales	2020-06-20
IT	2024-02-01

- 4) List the employee ids and last names of all employees who work in the HR department.

**Expected result:**

empid	Iname
1004	Baker
1005	Patel

- 5) List the project names and locations for all projects controlled by the Engineering department.

**Expected result:**

pname	plocation
Product Alpha	Cambridge

# *continued...*

**6) List the employee ids and project numbers for all employees who work more than 15 hours on a single project.**

**Expected result:**

empid	pno
1008	2002
1001	2004
1003	2004

**7) For each department, list the department name and the number of employees in that department.**

**Expected result:**

dname	num_employees
Engineering	3
Finance	1
HR	2
IT	1
Sales	1

# *continued...*

**8) List the employee ids and last names of employees who work on at least two projects.**

**Expected result:**

empid	Iname
1008	Ali

**9) List the employee ids and last names of employees whose salary is higher than the average salary in their own department.**

**Expected result:**

empid	Iname
1001	Rahman
1004	Baker

**10) List the employee ids and last names of employees who work on the maximum number of projects in the entire company.**

**Expected result:**

empid	Iname
1008	Ali

# More Queries to Attempt

11. *List the first and last names of all employees along with their department names.*
12. *List the project names and the last names of employees working on each project.*
13. *List the employee ids and total number of hours each employee works across all projects.*
14. *List the department names and the total salary paid to employees in each department.*
15. *List the ids and last names of employees who have no dependents.*
16. *List the project names of projects that currently have no employees assigned.*
17. *List the ids and last names of employees who work on every project.*
18. *List the department names that have more than three employees.*
19. *List the ids and last names of employees who earn the highest salary in their department.*
20. *List the ids and last names of employees whose total project hours exceed the company-wide average total project hours per employee.*